
Competing with SpaceX

Applied Data Science Capstone

Abhishek Deodhar

JULY 24, 2022

Table of Contents

- Executive summary
- Introduction
- Methodology
- Results & Discussion
- Conclusion
- Appendix

Executive summary

Main targets to be achieved :

- Collecting data from API and web scraping.
- Data wrangling.
- Exploratory Data Analysis with SQL.
- Exploratory Data Analysis with Data visualization.
- Interactive maps using Folium.
- Creating a Plotly Dash dashboard.
- Predictions using Machine learning.

Introduction

Background and Goal: SpaceX claims that it will cost them 62 million dollars as compared to other providers having a cost of more than 165M dollars. The main reason for this cost saving is reuse of the first stage.

Our goal in this project is to determine the cost of project by determining whether the first stage landing will be successful or not. We will achieve this by using machine learning techniques.

Data collection methodology

Objectives:

- Request to the SpaceX API
- Clean the requested data

Helper functions:

```
In [2]: # Takes the dataset and uses the rocket column to call the API and append the data to the List
def getBoosterVersion(data):
    for x in data['rocket']:
        response = requests.get("https://api.spacexdata.com/v4/rockets/"+str(x)).json()
        BoosterVersion.append(response['name'])
```

From the launchpad we would like to know the name of the launch site being used, the longitude, and the latitude.

```
In [3]: # Takes the dataset and uses the Launchpad column to call the API and append the data to the List
def getLaunchSite(data):
    for x in data['launchpad']:
        response = requests.get("https://api.spacexdata.com/v4/launchpads/"+str(x)).json()
        Longitude.append(response['longitude'])
        Latitude.append(response['latitude'])
        LaunchSite.append(response['name'])
```

Steps:

- Request APIs and normalize using json_normalize.
- Create helper functions for extracting info from API like boosterversion, Launchsite, Payload data etc.
- Apply the functions and create the data frame.
- Filter DF to include only falcon 9.
- Deal with the missing values.
- Refer Appendix slide no. 17,18.

Data Wrangling

1. As the first step we calculate the percentage of missing values & identified whether the columns are categorical or numerical.

```
In [6]: df.isnull().sum()/df.count()*100
```

```
FlightNumber    0.000
Date            0.000
BoosterVersion  0.000
PayloadMass     0.000
Orbit           0.000
LaunchSite      0.000
Outcome         0.000
Flights         0.000
GridFins        0.000
Reused          0.000
Legs            0.000
LandingPad      40.625
Block           0.000
ReusedCount     0.000
Serial          0.000
Longitude       0.000
Latitude        0.000
dtype: float64
```

```
In [7]: df.dtypes
```

```
FlightNumber    int64
Date            object
BoosterVersion  object
PayloadMass     float64
Orbit           object
LaunchSite      object
Outcome         object
Flights         int64
GridFins        bool
Reused          bool
Legs            bool
LandingPad      object
Block           float64
ReusedCount     int64
Serial          object
Longitude       float64
Latitude        float64
dtype: object
```

2. Then we calculate number of launches in each site.

```
In [8]: # Apply value_counts() on column LaunchSite
df['LaunchSite'].value_counts()
```

```
CCAFS SLC 40    55
KSC LC 39A      22
VAFB SLC 4E     13
Name: LaunchSite, dtype: int64
```

Data Wrangling

3. Calculate the number and occurrence of each orbit.
4. Count the different type of landing outcomes
5. Create a label where the value was 1 for successful landings and 0 for failures.
6. Store the landing values in a column called "class".
7. Finally determine the success rate by calculating the mean of the column class.

```
In [14]: # landing_class = 0 if bad_outcome
          # landing_class = 1 otherwise
          #df.head()

          def funA(input1):
              if input1 in bad_outcomes:
                  return 0
              else:
                  return 1

          landing_class= df["Outcome"].apply(funA)
          landing_class
```

```
0    0
1    0
2    0
3    0
4    0
..
85   1
86   1
87   1
88   1
89   1
```

```
Name: Outcome, Length: 90, dtype: int64
```

Data Wrangling

```
In [16]: df.head(5)
```

| LaunchSite | Outcome | Flights | GridFins | Reused | Legs | LandingPad | Block | ReusedCount | Serial | Longitude | Latitude | Class |
|--------------|----------------|---------|----------|--------|-------|------------|-------|-------------|--------|-------------|-----------|-------|
| CCAFS SLC 40 | None None | 1 | False | False | False | NaN | 1.0 | 0 | B0003 | -80.577366 | 28.561857 | 0 |
| CCAFS SLC 40 | None None | 1 | False | False | False | NaN | 1.0 | 0 | B0005 | -80.577366 | 28.561857 | 0 |
| CCAFS SLC 40 | None None | 1 | False | False | False | NaN | 1.0 | 0 | B0007 | -80.577366 | 28.561857 | 0 |
| VAFB SLC 4E | False Ocean | 1 | False | False | False | NaN | 1.0 | 0 | B1003 | -120.610829 | 34.632093 | 0 |
| CCAFS SLC 40 | None None | 1 | False | False | False | NaN | 1.0 | 0 | B1004 | -80.577366 | 28.561857 | 0 |

We can use the following line of code to determine the success rate:


```
In [17]: df["Class"].mean()
```

```
0.6666666666666666
```



EDA and interactive visual analytics methodology

Here we first load the data set and plot relevant parameters along the X and Y axis for exploring the data. The parameters explored were - FlightNumber vs. PayloadMass, FlightNumber vs LaunchSite, PayloadMass vs LaunchSite, Orbit vs Class.mean(),FlightNumber vs orbit, PayloadMass vs orbit, Year vs average success rate etc.

```
sns.catplot(y="LaunchSite", x="FlightNumber", hue="Class", data=df, aspect = 3)  
plt.xlabel("Flight Number",fontsize=20)  
plt.ylabel("Launch Site",fontsize=20)  
plt.show()
```



```
In [8]: # Plot a scatter point chart with x axis to be Payload and y axis to be the Orbit, and hue to be the class  
sns.catplot(y="Orbit", x="PayloadMass", hue="Class", data=df, aspect = 3)  
plt.xlabel("PayloadMass",fontsize=20)  
plt.ylabel("Orbit",fontsize=20)  
plt.show()
```



EDA and interactive visual analytics methodology

Feature engineering: Identifying the importance of each variable in finding the success rate. We apply one hot encoding to selected features - Orbits, LaunchSite, LandingPad, and Serial. This will make our data more relevant as we convert the categorical columns to a binary 0/1.

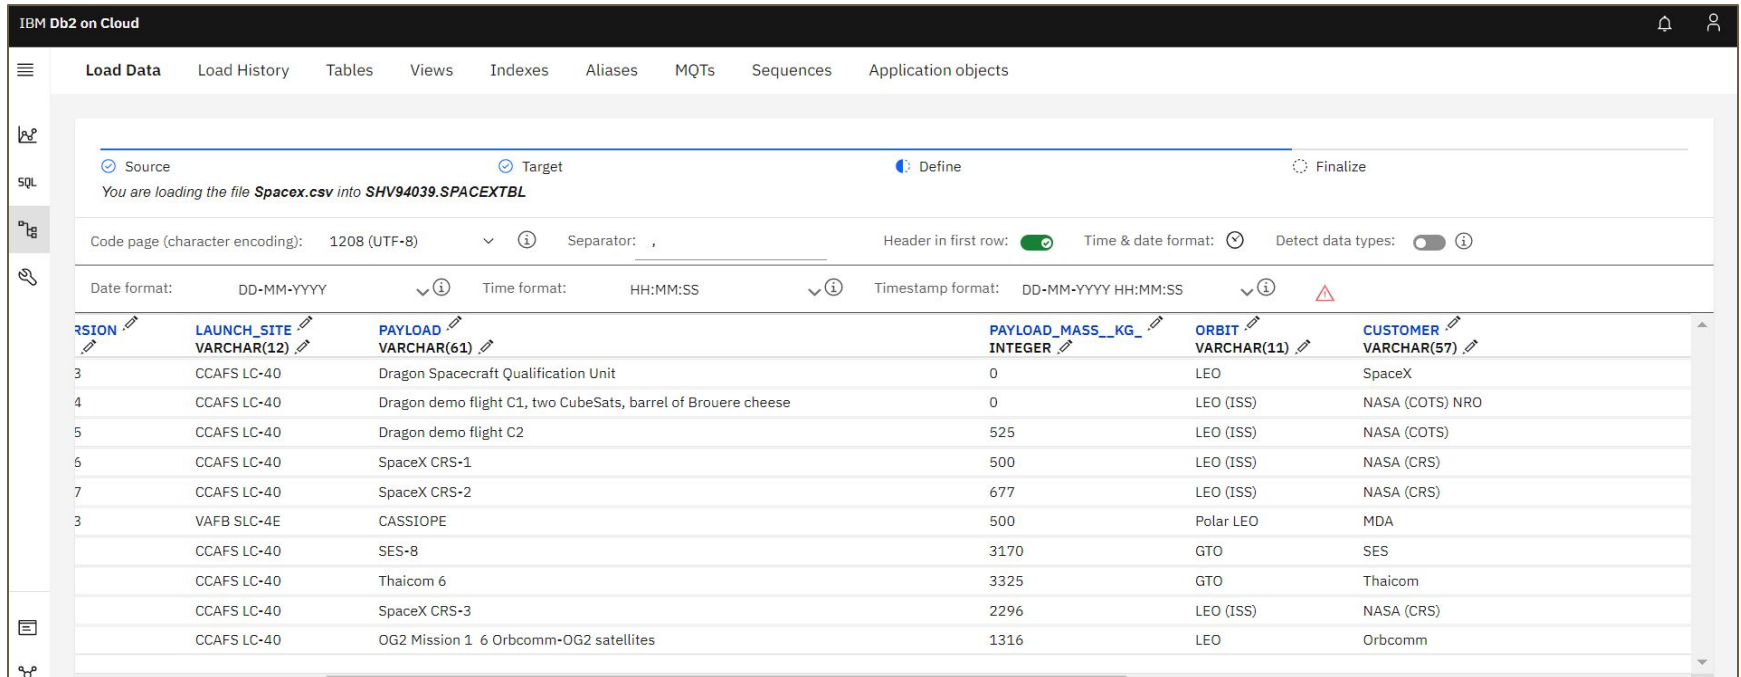
```
In [23]: # HINT: Use get_dummies() function on the categorical columns
features_one_hot = pd.get_dummies(features, columns=["Orbit", "LaunchSite", "LandingPad", "Serial"])
features_one_hot
```

| | FlightNumber | PayloadMass | Flights | GridFins | Reused | Legs | Block | ReusedCount | Orbit_ES-L1 | Orbit_GEO | ... | Serial_B104 |
|-----|--------------|--------------|---------|----------|--------|-------|-------|-------------|-------------|-----------|-----|-------------|
| 0 | 1 | 6104.959412 | 1 | False | False | False | 1.0 | 0 | 0 | 0 | ... | 0 |
| 1 | 2 | 525.000000 | 1 | False | False | False | 1.0 | 0 | 0 | 0 | ... | 0 |
| 2 | 3 | 677.000000 | 1 | False | False | False | 1.0 | 0 | 0 | 0 | ... | 0 |
| 3 | 4 | 500.000000 | 1 | False | False | False | 1.0 | 0 | 0 | 0 | ... | 0 |
| 4 | 5 | 3170.000000 | 1 | False | False | False | 1.0 | 0 | 0 | 0 | ... | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 85 | 86 | 15400.000000 | 2 | True | True | True | 5.0 | 2 | 0 | 0 | ... | 0 |
| 86 | 87 | 15400.000000 | 3 | True | True | True | 5.0 | 2 | 0 | 0 | ... | 0 |
| 87 | 88 | 15400.000000 | 6 | True | True | True | 5.0 | 5 | 0 | 0 | ... | 0 |
| 88 | 89 | 15400.000000 | 3 | True | True | True | 5.0 | 2 | 0 | 0 | ... | 0 |
| 89 | 90 | 3681.000000 | 1 | True | False | True | 5.0 | 0 | 0 | 0 | ... | 0 |

90 rows × 80 columns

EDA with SQL

Under this activity we load the data into the database, query the data using python, understand the data to get more information on parameters like launch sites, mission outcomes etc. Some of the tasks have been shown in the below slides. Refer appendix slide no. 19,20.



IBM Db2 on Cloud

Load Data | Load History | Tables | Views | Indexes | Aliases | MQTs | Sequences | Application objects

Source | Target | Define | Finalize

You are loading the file **Spacex.csv** into **SHV94039.SPACEXTBL**

Code page (character encoding): 1208 (UTF-8) | Separator: , | Header in first row: ☒ | Time & date format: | Detect data types: ☐

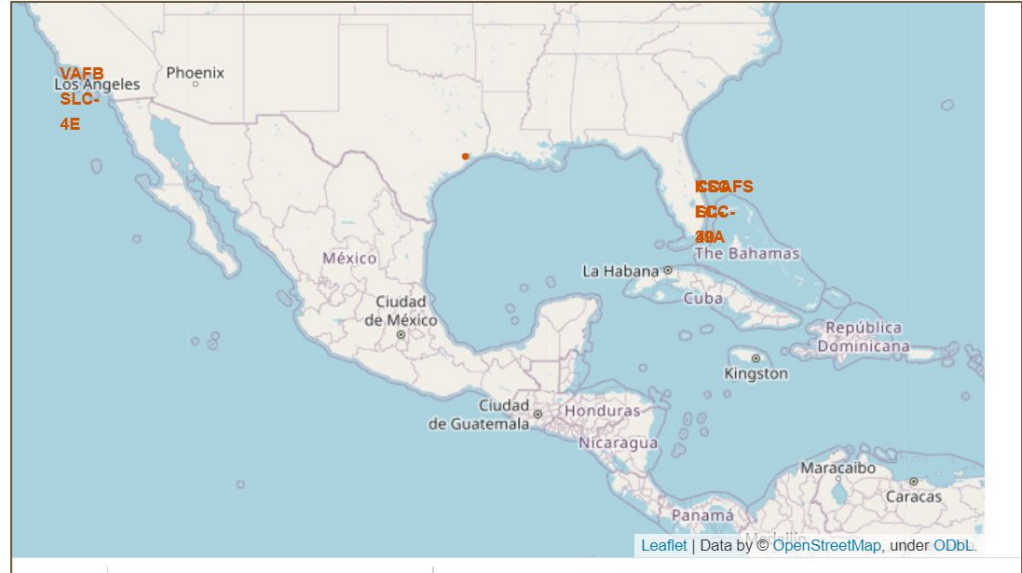
Date format: DD-MM-YYYY | Time format: HH:MM:SS | Timestamp format: DD-MM-YYYY HH:MM:SS

| MISSION | LAUNCH_SITE | PAYLOAD | PAYLOAD_MASS_KG | ORBIT | CUSTOMER |
|---------|-------------|---|-----------------|-----------|-----------------|
| 3 | CCAFS LC-40 | Dragon Spacecraft Qualification Unit | 0 | LEO | SpaceX |
| 4 | CCAFS LC-40 | Dragon demo flight C1, two CubeSats, barrel of Brouere cheese | 0 | LEO (ISS) | NASA (COTS) NRO |
| 5 | CCAFS LC-40 | Dragon demo flight C2 | 525 | LEO (ISS) | NASA (COTS) |
| 6 | CCAFS LC-40 | SpaceX CRS-1 | 500 | LEO (ISS) | NASA (CRS) |
| 7 | CCAFS LC-40 | SpaceX CRS-2 | 677 | LEO (ISS) | NASA (CRS) |
| 3 | VAFB SLC-4E | CASSIOPE | 500 | Polar LEO | MDA |
| | CCAFS LC-40 | SES-8 | 3170 | GTO | SES |
| | CCAFS LC-40 | Thaicom 6 | 3325 | GTO | Thaicom |
| | CCAFS LC-40 | SpaceX CRS-3 | 2296 | LEO (ISS) | NASA (CRS) |
| | CCAFS LC-40 | OG2 Mission 1 6 Orbcomm-OG2 satellites | 1316 | LEO | Orbcomm |

Interactive map with Folium

Here we are trying to plot launch sites and their proximity to key locations like the highway, coast etc. This helps us in understanding the logic behind selection of these sites at specific locations.


For e.g. most of the launch sites are near the coast and close to the equator. I think the main reason for that could be to get a radial path along the earth while exiting into outer space.





We also observed that the launch sites are close to railway routes and highways. It could be to facilitate easy transportation of large equipments.

Plotly Dash dashboard results

SpaceX Launch Records Dashboard



SpaceX Launch Records Dashboard



Predictive analysis (classification) results

The main purpose of this exercise is to predict whether first stage of falcon9 will land successfully or not. First we pre-process and standardize the dataset. Then we split the data into test-train samples. Post that we test the performance of Logistic Regression, Support Vector machines, Decision Tree Classifier, and K-nearest neighbors algorithms by training the models, performing grid-search and identifying the one with the best accuracy.

| | Model_Name | Score |
|---|--------------------|----------|
| 0 | LogisticRegression | 0.846429 |
| 1 | SVM | 0.848214 |
| 2 | DecisionTree | 0.887500 |
| 3 | KNN | 0.848214 |

```
In [120]: best=model_df[model_df['Score']==max(model_df['Score'])]
          best
          #print("The Best performing method is ",best['Model_Name'].apply(str)," with a score of ",best['Score'])
```

| | Model_Name | Score |
|---|--------------|--------|
| 2 | DecisionTree | 0.8875 |

Conclusions

- The main objective of this project was to identify whether a landing cheaper than what SpaceX is claiming is possible or not. Also, the focus was on the entire data science life cycle right from the retrieval of data till giving a result out. We started initially by understanding how to request data from an API.
- Then the next step was to wrangle the data, clean it and make it standard and easy to use. Post that we did Exploratory data analysis to understand the significance of each variable in the data set.
- Since this data set had some geographical relevance we used Folium to map certain coordinates and understand their relevance with the nearby geography. A plotly dashboard was also created to easily visualize and monitor the results.
- At the end we implemented a few classification algorithms to identify the success rate of the first landing. Among the various methods of classification the best one was identified.

Appendix

Appendix - Data collection methodology

- Request and parse the SpaceX launch data using the GET request

```
In [9]: static_json_url='https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork  
response1 = requests.get(static_json_url)
```

We should see that the request was successful with the 200 status response code

```
In [10]: response1.status_code
```

200

Now we decode the response content as a json using `.json()` and turn it into a Pandas dataframe using `json_normalize()`

```
In [11]: # Use json_normalize method to convert the json result into a dataframe  
df1 = response1.json()  
df2 = pd.json_normalize(df1)
```

Using the dataframe data print the first 5 rows

```
In [12]: # Get the head of the dataframe  
df2.head()
```

| | static_fire_date_utc | static_fire_date_unix | tbd | net | window | rocket | success | details | crew | ships | c |
|---|--------------------------|-----------------------|-------|-------|--------|--------------------------|---------|--|------|-------|----|
| 0 | 2006-03-17T00:00:00.000Z | 1.142554e+09 | False | False | 0.0 | 5e9d0d95eda69955f709d1eb | False | Engine failure at 33 seconds and loss of vehicle | [] | [] | [] |

Appendix - Data collection methodology

Filtering the data frame to only include Falcon 9 launches:

```
In [28]: # Hint data['BoosterVersion']!= 'Falcon 1'
data_falcon9=dataframe1[dataframe1['BoosterVersion']!='Falcon 1']
```

Now that we have removed some values we should reset the FlightNumber column

```
In [30]: data_falcon9.loc[:, 'FlightNumber'] = list(range(1, data_falcon9.shape[0]+1))
data_falcon9
```

C:\Users\User\anaconda3\lib\site-packages\pandas\core\indexing.py:1676: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
self._setitem_single_column(ilocs[0], value, pi)

| | FlightNumber | Date | BoosterVersion | PayloadMass | Orbit | LaunchSite | Outcome | Flights | GridFins | Reused | Legs |
|---|--------------|------------|----------------|-------------|-------|--------------|----------------|---------|----------|--------|---------------|
| 4 | 1 | 2010-06-04 | Falcon 9 | NaN | LEO | CCSFS SLC 40 | None None | 1 | False | False | False None |
| 5 | 2 | 2012-05-22 | Falcon 9 | 525.0 | LEO | CCSFS SLC 40 | None None | 1 | False | False | False None |
| 6 | 3 | 2013-03-01 | Falcon 9 | 677.0 | ISS | CCSFS SLC 40 | None None | 1 | False | False | False None |
| 7 | 4 | 2013-09-29 | Falcon 9 | 500.0 | PO | VAFB SLC 4E | False Ocean | 1 | False | False | False None |

Appendix - EDA with SQL

Task 1

Display the names of the unique launch sites in the space mission

```
In [8]: %sql select DISTINCT LAUNCH_SITE from SPACEXTBL
```

* sqlite:///my_data1.db
Done.

| Launch_Site |
|--------------|
| CCAFS LC-40 |
| VAFB SLC-4E |
| KSC LC-39A |
| CCAFS SLC-40 |

Task 2

Display 5 records where launch sites begin with the string 'CCA'

```
In [9]: %sql select * from SPACEXTBL where launch_site like 'CCA%' limit 5
```

* sqlite:///my_data1.db
Done.

| Date | Time (UTC) | Booster_Version | Launch_Site | Payload | PAYLOAD_MASS_KG | Orbit | Customer | Mission_Outcome | Landing_Outcome |
|------------|------------|-----------------|-------------|--|-----------------|-----------|-----------------|-----------------|---------------------|
| 04-06-2010 | 18:45:00 | F9 v1.0 B0003 | CCAFS LC-40 | Dragon Spacecraft Qualification Unit | 0 | LEO | SpaceX | Success | Failure (parachute) |
| 08-12-2010 | 15:43:00 | F9 v1.0 B0004 | CCAFS LC-40 | Dragon demo flight C1, two CubeSats, barrel of Brouere | 0 | LEO (ISS) | NASA (COTS) NRO | Success | Failure (parachute) |

Task 3

Display the total payload mass carried by boosters launched by NASA (CRS)

```
In [11]: %sql select sum(payload_mass_kg_) as sum from SPACEXTBL where customer like 'NASA (CRS)'
```

* sqlite:///my_data1.db
Done.

| sum |
|-------|
| 45596 |

Task 4

Display average payload mass carried by booster version F9 v1.1

```
In [12]: %sql select avg(payload_mass_kg_) as Average from SPACEXTBL where booster_version like 'F9 v1.1'
```

* sqlite:///my_data1.db
Done.

| Average |
|--------------------|
| 2534.6666666666665 |

Appendix - EDA with SQL

Task 5

List the date when the first succesful landing outcome in ground pad was acheived.

Hint: Use min function

```
In [72]: %sql select min(date) as Date from SPACEXTBL where mission_outcome like 'Success'
```

```
* sqlite:///my_data1.db
Done.
```

| Date |
|------------|
| 01-03-2013 |

Task 6

List the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000

```
In [99]: #%sql PRAGMA table_info(SPACEXTBL)

%sql select booster_version from SPACEXTBL where Landing_Outcome like 'Success (drone ship)' \
AND PAYLOAD_MASS_KG>4000 AND PAYLOAD_MASS_KG<6000
```

```
* sqlite:///my_data1.db
Done.
```

| Booster_Version |
|-----------------|
| F9 FT B1022 |
| F9 FT B1026 |
| F9 FT B1021.2 |
| F9 FT B1031.2 |

Task 7

List the total number of successful and failure mission outcomes

```
In [117]: %sql SELECT mission_outcome, count(*) as Count FROM SPACEXTBL GROUP by\
mission_outcome ORDER BY mission_outcome
```

```
* sqlite:///my_data1.db
Done.
```

| Mission_Outcome | Count |
|----------------------------------|-------|
| Failure (in flight) | 1 |
| Success | 98 |
| Success | 1 |
| Success (payload status unclear) | 1 |

Task 8

List the names of the booster_versions which have carried the maximum payload mass. Use a subquery

```
In [105]: %sql select booster_version from SPACEXTBL where \
PAYLOAD_MASS_KG=(select max(PAYLOAD_MASS_KG) from SPACEXTBL)
```

```
* sqlite:///my_data1.db
Done.
```

| Booster_Version |
|-----------------|
| F9 B5 B1048.4 |
| F9 B5 B1049.4 |
| F9 B5 B1051.3 |
| F9 B5 B1056.4 |

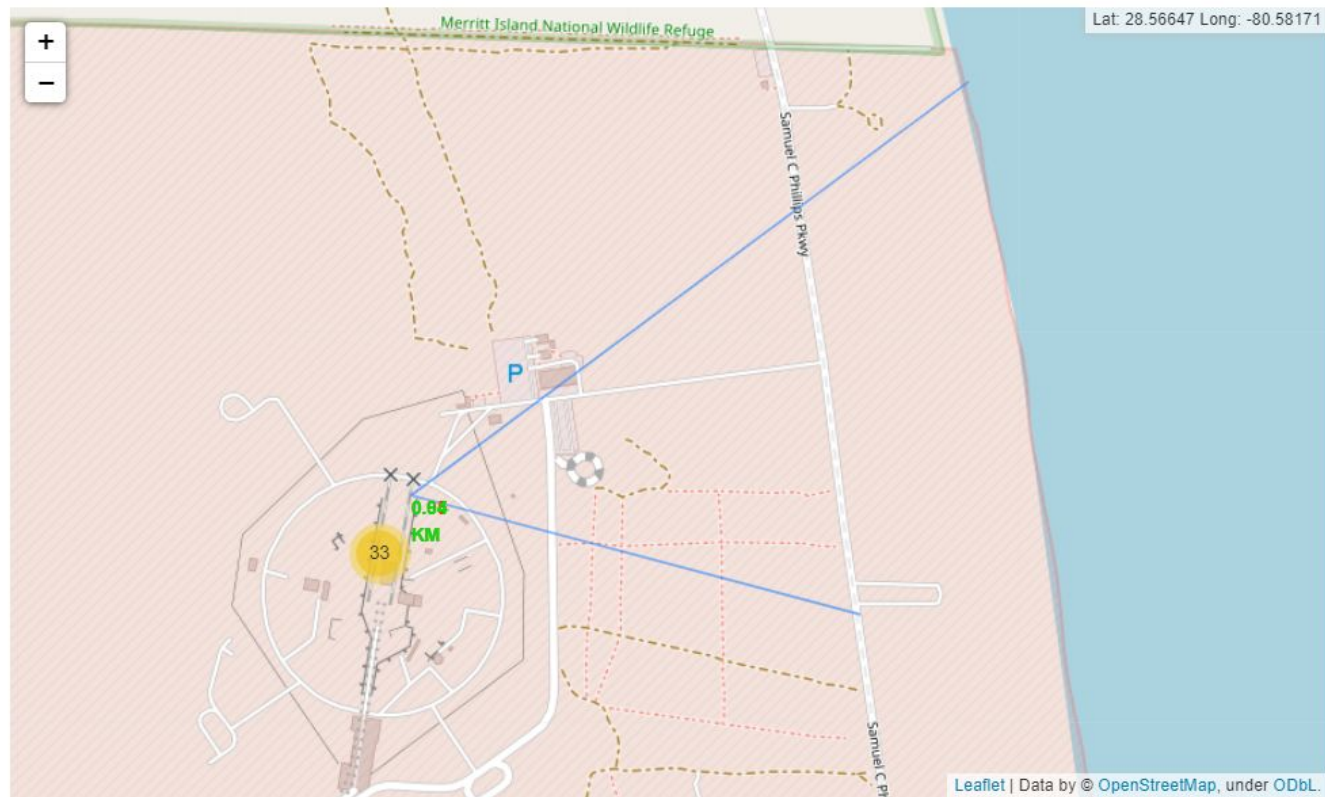
site_map_new

A map of Merritt Island National Wildlife Refuge. A yellow circle with the number '33' inside is located in the lower-left quadrant. A blue line extends from the center of this circle towards the upper-right, ending near a body of water. The text '0.96 KM' is written in black above the yellow circle. The map shows various land parcels, some with dashed yellow boundaries, and a road labeled 'Samuel C. Ripley Road'. The text 'Merritt Island National Wildlife Refuge' is at the top. In the bottom right corner, it says 'Leaflet | Data by © OpenStreetMap, under ODbL'. Coordinates 'Lat: 28.56489 Long: -80.56763' are in the top right corner. A zoom control is in the top left.

TODO: For each launch result in `spacex_df` data frame, add a `folium.Marker` to `marker_cluster`

Appendix - Interactive map with Folium

```
lines2=folium.PolyLine(locations=[[launch_site_lat,launch_site_lon],[highway_lat,highway_lon]], weight=1)
site_map_new.add_child(lines2)
```



Appendix - Predictive analysis

TASK 1

Create a NumPy array from the column `class` in `data`, by applying the method `to_numpy()` then assign it to the variable `Y`, make sure the output is a Pandas series (only one bracket `df['name of column']`).

```
In [5]: Y=pd.DataFrame.to_numpy(data['Class'])
```

```
Y
array([[0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1,
        1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1], dtype=int64)
```

TASK 3

Use the function `train_test_split` to split the data `X` and `Y` into training and test data. Set the parameter `test_size` to 0.2 and `random_state` to 2. The training data and test data should be assigned to the following labels.

`X_train, X_test, Y_train, Y_test`

```
In [11]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
print ('Train set:', X_train.shape, y_train.shape)
print ('Test set:', X_test.shape, y_test.shape)
```

```
Train set: (72, 83) (72,)
Test set: (18, 83) (18,)
```

TASK 2

Standardize the data in `x` then reassign it to the variable `x` using the transform provided below.

```
In [10]: X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]
```

```
array([[ -1.71291154e+00,  -3.32153339e-17,  -6.53912840e-01,
        -1.57589457e+00,  -9.73440458e-01,  -1.05999788e-01,
        -1.05999788e-01,  -6.54653671e-01,  -1.05999788e-01,
        -5.51677284e-01,   3.44342023e+00,  -1.85695338e-01,
        -3.33333333e-01,  -1.05999788e-01,  -2.42535625e-01,
        -4.29197538e-01,   7.97724035e-01,  -5.68796459e-01,
        -4.10890702e-01,  -4.10890702e-01,  -1.50755672e-01,
```

TASK 4

Create a logistic regression object then create a `GridSearchCV` object `logreg_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary parameters .

```
In [13]: parameters = {'C':[0.01,0.1,1],
                       'penalty':['l2'],
                       'solver':['lbfgs']}

In [14]: parameters = {'C':[0.01,0.1,1], 'penalty':['l2'], 'solver':['lbfgs']}# L1 Lasso L2 ridge
#parameters = [{'max_depth': List(range(10, 15)), 'max_features': List(range(0,14))}]

lr=LogisticRegression()

logreg_cv = GridSearchCV(lr, parameters, cv=10,scoring='accuracy')

logreg_cv.fit(X_train, y_train)

logreg_cv

GridSearchCV(cv=10, estimator=LogisticRegression(),
             param_grid={'C': [0.01, 0.1, 1], 'penalty': ['l2'],
                         'solver': ['lbfgs']},
             scoring='accuracy')
```

Appendix - Predictive analysis

TASK 5


Calculate the accuracy on the test data using the method `score` :

```
In [16]: #Logreg_cv.fit(X_test, y_test)
#Logreg_cv.fit(X_train, y_train)
print("accuracy :", logreg_cv.score(X_test, y_test))
```

```
accuracy : 0.8333333333333334
```

Lets look at the confusion matrix:

```
In [17]: yhat=logreg_cv.predict(X_test)
plot_confusion_matrix(y_test,yhat)
```



TASK 7


Calculate the accuracy on the test data using the method `score` :

```
In [21]: #svm_cv.fit(X_test, y_test)
print("accuracy :", svm_cv.score(X_test, y_test))
```

```
accuracy : 0.8333333333333334
```

We can plot the confusion matrix

```
In [22]: yhat=svm_cv.predict(X_test)
plot_confusion_matrix(y_test,yhat)
```



TASK 6

Create a support vector machine object then create a `GridSearchCV` object `svm_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters` .

```
In [18]: parameters = {'kernel':('linear', 'rbf','poly','rbf', 'sigmoid'),
                      'C': np.logspace(-3, 3, 5),
                      'gamma':np.logspace(-3, 3, 5)}
svm = SVC()
```

```
In [19]: svm_cv = GridSearchCV(svm, parameters, cv=10, scoring='accuracy')
svm_cv.fit(X_train, y_train)
```

```
GridSearchCV(cv=10, estimator=SVC(),
             param_grid={'C': array([1.00000000e-03, 3.16227766e-02, 1.00000000e+00, 3.16227766e+01,
1.00000000e+03]),
                        'gamma': array([1.00000000e-03, 3.16227766e-02, 1.00000000e+00, 3.16227766e+01,
1.00000000e+03]),
                        'kernel': ('linear', 'rbf', 'poly', 'rbf', 'sigmoid')},
             scoring='accuracy')
```

```
In [20]: print("tuned hyperparameters :(best parameters) ",svm_cv.best_params_)
print("accuracy :",svm_cv.best_score_)
```

```
tuned hyperparameters :(best parameters) {'C': 1.0, 'gamma': 0.03162277660168379, 'kernel': 'sigmoid'}
accuracy : 0.8482142857142856
```

TASK 8

Create a decision tree classifier object then create a `GridSearchCV` object `tree_cv` with `cv = 10`. Fit the object to find the best parameters from the dictionary `parameters` .

```
In [23]: parameters = {'criterion': ['gini', 'entropy'],
                      'splitter': ['best', 'random'],
                      'max_depth': [2*n for n in range(1,10)],
                      'max_features': ['auto', 'sqrt'],
                      'min_samples_leaf': [1, 2, 4],
                      'min_samples_split': [2, 5, 10]}
```

```
tree = DecisionTreeClassifier()
```

```
In [24]: tree_cv = GridSearchCV(tree, parameters, cv=10, scoring='accuracy')
tree_cv.fit(X_train, y_train)
```

```
GridSearchCV(cv=10, estimator=DecisionTreeClassifier(),
             param_grid={'criterion': ['gini', 'entropy'],
                        'max_depth': [2, 4, 6, 8, 10, 12, 14, 16, 18],
                        'max_features': ['auto', 'sqrt'],
                        'min_samples_leaf': [1, 2, 4],
                        'min_samples_split': [2, 5, 10],
                        'splitter': ['best', 'random']},
             scoring='accuracy')
```